

To process the original list of 1000 tweets, we used a series of five filters and received an output of 516 tweets. Of these tweets, we used a method borrowed from the original paper to search for literariness. It counted 455 literary tweets and 61 non-literary tweets, a ratio of 13%- a far cry from the full list's 37% ratio. The primary reason for this is that there were a massive number of bot tweets in the initial input that are not filtered out for being duplicates, because they all tag a different account. These tweets are labeled as literary by the function and no doubt skew the percentage by their sheer volume.

On top of literariness we also evaluated some example tweets by several criteria: RC-type, RC-head, and Role. The results of that are as follows:

RC-Type	RC-Head	Role
Restrictive (128)	You (149)	Other (149)
Appositive (22)	Her (1)	Obj (1)

That data was gathered from a slice of our larger table, literary-annotated-tweets.csv.

The data on that is as follows:

"he who"	0
"she who"	0
"it who"	0
"him who"	0
"her who"	1
"they who"	0
"them who"	0
"we who"	0
"us who"	3
"you who"	511

While the numbers look incredibly skewed, the filter is actually working as intended. Our original list of 1000 contained mostly *you who* constructs, which is reflected in our output.

Despite the filters working as intended, we do not believe our data is comparable to CTK's. While we did notice some anomalous tweets (mostly bot tweets) in the final filtered list, there is simply too much discrepancy between our results and CTK's to call this an accurate re-creation. For instance our number of Restrictives compared to CTK's was far too high to be brushed off as resulting from abnormality of the set.